

Data Selection via Optimal Control for Language Models

Yuxian Gu^{1,2}, Li Dong², Hongning Wang¹, Yaru Hao²,
Qingxiu Dong², Minlie Huang¹, Furu Wei²

¹The CoAI Group, Tsinghua University

²Microsoft Research





清华大学
Tsinghua University



Data challenges for pre-training LMs




- Large amount of data makes pre-training quite **inefficient**.



Data challenges for pre-training LMs



- High-quality pre-training data is running out.





Models consume faster than humans produce.

Possible Solution: Data Selection?




- Select a Pre-Training Corpus Subset for Better Target Performance
 - Target: Math, Code, High-Quality Instruction, etc..



Challenges for Data Selection



- Current data selection/filtering is heuristic-based and tricky task.



Overview



○ Data challenges for pre-training LMs


- ◆ Large amount of data makes pre-training quite **inefficient**.
- ◆ High-quality pre-training data is **running out**.
- ◆ Data selection/cleaning is a **heuristic-based** tricky task.

○ PDS: Data Selection via Optimal Control for Pre-Training


Theorem 2.1: PMP Conditions for Data Selection

$$\begin{aligned}\theta_{t+1}^* &= \theta_t^* - \eta \nabla L(\theta_t^*, \gamma^*), \quad \theta_0^* = \theta_0, \\ \lambda_t^* &= \lambda_{t+1}^* + \nabla J(\theta_t^*) - \eta \nabla^2 L(\theta_t^*, \gamma^*) \lambda_{t+1}^*, \\ \gamma^* &= \arg \max_{\gamma} \sum_{n=1}^{|\mathcal{D}|} \gamma_n \left[\sum_{t=0}^{T-1} \lambda_{t+1}^{*\top} \nabla l(x_n, \theta_t^*) \right]\end{aligned}$$

Good theoretical foundation



2x acceleration




Improvement on limited data

Formulate Data Selection





- Optimize the data selection strategy for lower downstream loss


γ : indicates a sample is selected or not



$$\gamma = [1, 1, 1, \dots, 1]$$



$$\gamma^* = [1, 0, 1, \dots, 0]$$




Training on the Selected Data



- Loss: Treat the γ as the weights of the instance losses:

$$L(\theta, \gamma) = \sum_{n=1}^{|D|} \gamma_n l(x_n, \theta)$$

$\gamma_n = 1$: $l(x_n, \theta)$ is selected
 $\gamma_n = 0$: $l(x_n, \theta)$ is ignored




- The model is trained with:

$$\theta_{t+1} = \theta_t - \eta \nabla L(\theta_t, \gamma)$$

Target to Optimize


- $J(\theta_t)$: downstream loss to minimize (like on math, code, etc.)
- Minimizing the Area Under the Loss curve (AUC)
 - ◆ AUC is directly related to the Scaling Law constants (see Appendix A in our paper)
 - ◆ Optimizing AUC is improving the Scaling Law!



$$\begin{aligned} & \min_{\gamma} \sum_{t=1}^T J(\theta_t), \\ \text{s.t. } & \theta_{t+1} = \theta_t - \eta \nabla L(\theta_t, \gamma) \end{aligned}$$

Target to Optimize

- $J(\theta_t)$: downstream loss to minimize (like on math, code, etc.)
- Minimizing the Area Under the Loss curve (AUC)
 - ◆ AUC is directly related to the Scaling Law constants (see Appendix A in our paper)
 - ◆ Optimizing AUC is improving the Scaling Law!



$$\begin{aligned} \min_{\gamma} \quad & \sum_{t=1}^T J(\theta_t), \\ \text{s.t. } \theta_{t+1} = & \theta_t - \eta \nabla L(\theta_t, \gamma) \end{aligned}$$


Hard to solve? Optimal Control!

Data Selection as a Control Problem




- Original Problem: Optimize the data selection strategy


γ : indicates a sample is selected or not



$$\gamma = [1, 1, 1, \dots, 1]$$



$$\gamma^* = [1, 0, 1, \dots, 0]$$




Data Selection as a Control Problem




- Analogy: Optimizing fuel use when flying a rocket
 - Data is the “fuel” in pre-training language models



Original Fuel Use




Optimized Fuel Use




Mathematical Equivalence to Optimal Control



Data Selection for LMs



Fuel Use Optimization



Control Variable	Selection Strategy: γ	Fuel Consumption: u_t
Objective	Minimal AUC: $\min_{\gamma} \sum_{t=1}^T J(\theta_t),$	Maximize Distance: $\max_{u_t} x = \sum_{t=0}^T v_t \Delta t$
Constraints	Regularity: $\gamma \in U.$	Constant Total Fuel: $\sum_{t=0}^T u_t = U$
Dynamics	Gradient decent: $\theta_{t+1} = \theta_t - \eta \nabla L(\theta_t, \gamma)$	Newton's Law: $\frac{v_{t+\Delta t} - v_t}{\Delta t} = -mg + F(u_t)$

Solving the Problem

○ Pontryagin's Maximum Principle (PMP)

◆ Gives necessary conditions for the **optimality** of the problem

$$\min_{\gamma} \sum_{t=1}^T J(\theta_t),$$

$$\text{s.t. } \theta_{t+1} = \theta_t - \eta \nabla L(\theta_t, \gamma)$$

$$L(\theta, \gamma) = \sum_{n=1}^{|\mathcal{D}|} \gamma_n l(x_n, \theta)$$



Lev Pontryagin, 1908 - 1988

Conditions for Optimal Data Selection



Theorem 2.1 (PMP Conditions for Data Selection).

$$\left\{ \begin{array}{l} \#1 \quad \theta_{t+1}^* = \theta_t^* - \eta \nabla L(\theta_t^*, \gamma^*), \\ \#2 \quad \lambda_t^* = \lambda_{t+1}^* + \nabla J(\theta_t^*) - \eta \nabla^2 L(\theta_t^*, \gamma^*) \lambda_{t+1}^*, \\ \#3 \quad \boxed{\gamma^*} = \arg \max_{\gamma} \sum_{n=1}^{|\mathcal{D}|} \gamma_n \left[\sum_{t=0}^{T-1} \lambda_{t+1}^{*\top} \nabla l(x_n, \theta_t^*) \right] \end{array} \right.$$

We can get the optimal data score γ^* here!


#1: Learning Condition

$$\boxed{\theta_{t+1}^* = \theta_t^* - \eta \nabla L(\theta_t^*, \gamma^*)}$$

Model Parameters

Optimal Data Selection Strategy

- Exactly the parameter updating policy of training LMs
- Constrains the θ_t^* to be reachable with GD under the optimal data selection



#2: Target Condition


$$\lambda_t^* = \lambda_{t+1}^* + \nabla J(\theta_t^*) - \eta \nabla^2 L(\theta_t^*, \gamma^*) \lambda_{t+1}^*$$

Ideal gradient = Target + Learning dynamics

- λ_t^* : the **ideal gradient** of the high-quality data points.

 **Compass** for high-quality data.


- Ideal gradient includes information of **Target Loss** and **Learning Dynamics**.



#3: Maximum Condition

$$\gamma^* = \arg \max_{\gamma} \sum_{n=1}^{|\mathcal{D}|} \gamma_n \left[\sum_{t=0}^{T-1} \lambda_{t+1}^{*\top} \nabla l(x_n, \theta_t^*) \right]$$

Gradient of each sample



$$\sum_t \lambda_{t+1}^{*\top} \nabla l_{n,t} < \sum_t \lambda_{t+1}^{*\top} \nabla l_{m,t} < \sum_t \lambda_{t+1}^{*\top} \nabla l_{k,t}$$

$$\Rightarrow \quad \gamma_n < \gamma_m < \gamma_k$$

Examples with **closer gradients to λ_t^*** should **have higher γ** .

Summing up

Theorem 2.1 (PMP Conditions for Data Selection).

#1 Learning Condition $\theta_{t+1}^* = \theta_t^* - \eta \nabla L(\theta_t^*, \gamma^*)$

#2 Target Condition $\lambda_t^* = \lambda_{t+1}^* + \nabla J(\theta_t^*) - \eta \nabla^2 L(\theta_t^*, \gamma^*) \lambda_{t+1}^*$

#3 Maximum Condition $\gamma^* = \arg \max_{\gamma} \sum_{n=1}^{|\mathcal{D}|} \gamma_n \left[\sum_{t=0}^{T-1} \lambda_{t+1}^{*\top} \nabla l(x_n, \theta_t^*) \right]$

How to Solve (Theoretically)?



Use learning condition to forward compute θ_0 to θ_T

forward Pass with #1: $\theta_0 \rightarrow \theta_1 \rightarrow \dots \rightarrow \theta_T$

#2 Target Condition $\lambda_t^* = \lambda_{t+1}^* + \nabla J(\theta_t^*) - \eta \nabla^2 L(\theta_t^*, \gamma^*) \lambda_{t+1}^*$

#3 Maximum Condition $\gamma^* = \arg \max_{\gamma} \sum_{n=1}^{|\mathcal{D}|} \gamma_n \left[\sum_{t=0}^{T-1} \lambda_{t+1}^{*\top} \nabla l(x_n, \theta_t^*) \right]$

How to Solve (Theoretically)?



Use Target Condition to reverse compute λ_T to λ_0

forward Pass with #1: $\theta_0 \rightarrow \theta_1 \rightarrow \dots \rightarrow \theta_T$

reverse Pass with #2: $\lambda_0 \leftarrow \lambda_1 \leftarrow \dots \leftarrow \lambda_T$

#3 Maximum Condition $\gamma^* = \arg \max_{\gamma} \sum_{n=1}^{|\mathcal{D}|} \gamma_n \left[\sum_{t=0}^{T-1} \lambda_{t+1}^{*\top} \nabla l(x_n, \theta_t^*) \right]$


How to Solve (Theoretically)?

Use Maximum Condition to solve the final γ

forward Pass with #1: $\theta_0 \rightarrow \theta_1 \rightarrow \dots \rightarrow \theta_T$

reverse Pass with #2: $\lambda_0 \leftarrow \lambda_1 \leftarrow \dots \leftarrow \lambda_T$

maximum γ with #3:



How to Solve (Theoretically)?


- Iteratively Solving γ until convergence (Algorithm 1)

while γ not converged, do

forward Pass with #1: $\theta_0 \rightarrow \theta_1 \rightarrow \dots \rightarrow \theta_T$

reverse Pass with #2: $\lambda_0 \leftarrow \lambda_1 \leftarrow \dots \leftarrow \lambda_T$



maximum γ with #3:



Equivalence to “Meta” Gradient Decent



- Optimizing the whole training process with “Meta GD”
- A training process can be viewed as an NN (vertically)



Solve γ with “Meta” GD


- Iteratively Solving γ until convergence (Algorithm 1)

while γ not converged, do

forward Pass with #1: $\theta_0 \rightarrow \theta_1 \rightarrow \dots \rightarrow \theta_T$

reverse Pass with #2: $\lambda_0 \leftarrow \lambda_1 \leftarrow \dots \leftarrow \lambda_T$

maximum γ with #3:



Solve γ with “Meta” GD




- Iteratively Solving γ until convergence (Algorithm 1)

while γ not converged, do

Forward pass of “Meta” GD: $\theta_0 \rightarrow \theta_1 \rightarrow \dots \rightarrow \theta_T$

Backward pass of “Meta” GD: $\lambda_0 \leftarrow \lambda_1 \leftarrow \dots \leftarrow \lambda_T$

Step pass of “Meta” GD:



How to Solve (Theoretically)?

- Iteratively Solving γ (Algorithm 1)



while γ not converged, do

forward Pass with #1: $\theta_0 \rightarrow \theta_1 \rightarrow \dots \rightarrow \theta_T$

reverse Pass with #2: $\lambda_0 \leftarrow \lambda_1 \leftarrow \dots \leftarrow \lambda_T$

maximum γ with #3:

$$\eta \nabla^2 L(\theta_t^*, \gamma^*) \lambda_{t+1}^*$$




Forward and Reverse passes are computationally intensive!

Efficient Implementation: Proxy to Large




- ① Solve γ on a small model (e.g., 140M) and data (e.g., 160M tokens)



Efficient Implementation: Proxy to Large




- ① Solve γ on a small model (e.g., 140M) and data (e.g., 160M tokens)
- ② Fit γ with a data scorer (e.g., a 140M LM with a regression head)



Efficient Implementation: Proxy to Large




- ① Solve γ on a small model (e.g., 140M) and data (e.g., 160M tokens)
- ② Fit γ with a data scorer (e.g., a 140M LM with a regression head)
- ③ Infer γ on the whole dataset (e.g., 100B tokens)



Efficient Implementation: Proxy to Large




- ① Solve γ on a small model (e.g., 140M) and data (e.g., 160M tokens)
- ② Fit γ with a data scorer (e.g., a 140M LM with a regression head)
- ③ Infer γ on the whole dataset (e.g., 100B tokens)



Efficient Implementation: Proxy to Large



PMP-based Data Selection (PDS) is Efficient




Experiment Setups



○ Training & Evaluation Setups

- ◆ Pre-training LMs from scratch
- ◆ Evaluate zero-shot performance



○ Data Setups


- ◆ Pre-Training data: CommonCrawl from Redpajama (100B tokens)
- ◆ Downstream loss $J(\theta)$: loss on LIMA (1k high-quality instruction-response pairs)
- ◆ Evaluation: 9 common NLP benchmarks: LAMBADA, Hellaswag, BoolQ, etc.

Performance Improvement



- Select 50B-token corpus from 125B-token corpus.
- Match the total training steps with the baselines (training computation)


	HS	LAMB	Wino.	OBQA	ARC-e	ARC-c	PIQA	SciQ	BoolQ	Avg.
Model Size = 470M										
Conventional	36.7	41.4	52.4	30.4	44.8	25.2	61.0	70.6	60.4	47.0
RHO-Loss	36.6	42.4	53.0	29.4	43.7	25.2	60.4	72.8	59.8	47.0
DSIR	36.4	42.6	51.7	29.8	46.0	24.7	61.0	72.0	55.8	46.7
IF-Score	36.6	41.8	53.4	29.6	44.7	25.1	60.8	68.8	58.7	46.6
PDS	37.9	44.6	52.3	29.8	46.5	25.8	61.8	73.8	61.4	48.2
Model Size = 1B										
Conventional	39.9	47.6	52.4	30.6	49.3	26.4	63.1	73.7	60.9	49.3
RHO-Loss	39.8	47.0	53.0	30.8	48.0	26.4	62.9	71.1	61.0	48.9
DSIR	40.8	47.8	53.0	31.2	49.8	26.8	62.7	76.6	58.0	49.6
IF-Score	39.4	47.0	52.6	28.6	49.4	26.4	63.5	74.0	60.5	49.0
PDS	42.1	48.8	54.0	33.4	51.3	28.0	64.1	78.5	58.7	51.0



Extrapolation to 400B models on 15T tokens



- Fitting the model loss curves with the Scaling Law



- Extrapolating to 400B models on 10T tokens


$$L(N, D) = E + \frac{A}{N^\alpha} + \frac{B}{D^\beta}$$

	N	D	Conventional	PDS
GPT-3 [12]	175B	300B	2.882	2.872
Llama [72]	6.7B	1.0T	2.942	2.896
Llama 2 [73]	70B	2.0T	2.877	2.855
Llama 3.1 [21]	405B	15T	2.851	2.838

Computation Saving



- 2.0x acceleration on 1.7B models



- PDS is efficient and offline



- Select data once for all models

	FLOPs ($\times 10^{20}$)	Actual Time
PDS	Proxy γ -solver	0.49 15.2 Hours
	Data Scorer	0.063 1.50 Hours
	Data Selection	0.0 10.2 Minutes
Pre-Training	5.1	144 Hours

Data Utilization Improvement







○ Performance improvement with limit data (50B tokens)



Conclusion

- A novel perspective for Data selection: Optimal Control problem



- ◆ Good theoretical guarantees ✓
- ◆ Efficient Implementation ✓
- ◆ Sound empirical results ✓

A **rigorous, theory-driven alternative** to the ad-hoc practices that currently dominate LM pre-training



Thanks!

-  Paper: <https://arxiv.org/abs/2410.07064>
-  GitHub: https://github.com/microsoft/LMOps/tree/main/data_selection
-  HuggingFace: <https://huggingface.co/Data-Selection>

Paper:



Code:



HF:

